

Stats and Public Health Part 2: Data Analysis

Daniel Mortensen

3-Aug-2021

Imports

Library imports

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

from scipy import stats
from statsmodels.stats.proportion import proportions_ztest
```

Importing the mosquito tracking data for the West Nile Virus.

```
In [2]: mosquito_original = pd.read_csv('./mosquito_data_part_2.csv')
mosquito_df = mosquito_original.copy()
```

```
In [3]: mosquito_df.shape
```

```
Out[3]: (18495, 12)
```

```
In [4]: mosquito_df.isna().sum()
```

```
Out[4]: Year                0
Week                0
Address Block        0
Trap                0
Trap type            0
Date                0
Mosquito number      0
WNV Present          0
Species              0
Lat                 0
Lon                 0
Month               0
dtype: int64
```

Data Management Plan

- Year - Numeric --> **no change**
- Week - Numeric --> could change to cyclic variable, but we don't need weekly granularity for this report --> **drop column**

- Address Block - Categorical --> data is repeated in `Lon` and `Lat` columns --> **drop column**
- Trap - Categorical --> Essentially synonymous with `Address Block` and `Lon` and `Lat` columns --> **drop column**
- Trap type - Categorical --> **Convert to dummy variable**
- Date - Numeric --> Already represented by `Year`, `Month` and `Week` columns --> **drop column**
- Mosquito number - Numeric --> **no change**
- WNV Present - Categorical --> **convert to binary**
- Species - Categorical --> **convert to dummy variables**
- Lat - Numeric --> **no change**
- Lon - Numeric --> **no change**
- Month - Numeric --> **change to cyclic variable**

```
In [5]: mosquito_df.drop(columns=["Week", "Address Block", "Trap", "Date"], inplace=True)
```

Determining column data types and splitting them into a numeric data frame and a categorical data frame for processing.

```
In [6]: print(mosquito_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18495 entries, 0 to 18494
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  18495 non-null  int64
1   Trap type             18495 non-null  object
2   Mosquito number       18495 non-null  int64
3   WNV Present           18495 non-null  object
4   Species               18495 non-null  object
5   Lat                   18495 non-null  float64
6   Lon                   18495 non-null  float64
7   Month                 18495 non-null  int64
dtypes: float64(2), int64(3), object(3)
memory usage: 1.1+ MB
None
```

```
In [7]: df_numeric = mosquito_df.select_dtypes(["int64", "float64", "uint8", "int32"])
df_categorical = mosquito_df.select_dtypes("object")
```

Instructions

Now that you are familiar with the data, we will move on to a set of analyses on the relationship between the different variables and the mosquito number, as well as the probability of finding West Nile Virus (WNV) at any particular time and location.

Part 1 - Basic Analysis

1. Convert the `WNV Present` column into a binary column, and create dummy variables from the `Trap type` column.

Determining possible values for `WNV Present`

```
In [8]: df_categorical["WNV Present"].value_counts()
```

```
Out[8]: negative    14501
        positive     3994
        Name: WNV Present, dtype: int64
```

Possible values include "negative" and "positive". Replacing "negative" with "0" and "positive" with "1".

```
In [9]: WNV_neg = (df_categorical['WNV Present'] == "negative")
        WNV_zero = (df_categorical['WNV Present'] == 0)

        df_numeric['WNV Present'] = np.where(WNV_neg | WNV_zero, 0, 1)
        mosquito_df['WNV Present'] = np.where(WNV_neg | WNV_zero, 0, 1)
        df_categorical.drop(columns="WNV Present", inplace=True)
        display(df_categorical.head(5))
```

<ipython-input-9-9f628e3d596f>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_numeric['WNV Present'] = np.where(WNV_neg | WNV_zero, 0, 1)
C:\Users\Daniel\anaconda3\lib\site-packages\pandas\core\frame.py:4308: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
return super().drop()
```

	Trap type	Species
0	GRAVID	CULEX RESTUANS
1	GRAVID	CULEX RESTUANS
2	GRAVID	CULEX RESTUANS
3	GRAVID	CULEX RESTUANS
4	GRAVID	CULEX RESTUANS

Creating dummy variables for trap type. There is not a clear best trap type to drop, so I will simply drop the first one, which is `CDC`

```
In [10]: trap_type_dummies = pd.get_dummies(df_categorical[["Trap type"]], drop_first=True)

        df_numeric = pd.concat([df_numeric, trap_type_dummies], axis = 1)
        df_numeric.head(3)
        df_categorical.drop(columns="Trap type", inplace=True)
```

C:\Users\Daniel\anaconda3\lib\site-packages\pandas\core\frame.py:4308: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
return super().drop()
```

Determining possible Species and their counts

```
In [11]: df_categorical.value_counts(normalize=True)*100
```

```
Out[11]: Species
CULEX RESTUANS      64.157881
CULEX PIPIENS       29.662071
CULEX TERRITANS      4.958097
CULEX SALINARIUS     1.221952
dtype: float64
```

Changing `Species` to a dummy variable. Species CULEX SALINARIUS is the least populous, representing only about 1% of all mosquitos caught, so I will drop the dummy variable for this species.

```
In [12]: species_dummies = pd.get_dummies(df_categorical[["Species"]]).drop(columns="Species_CUL
df_numeric = pd.concat([df_numeric, species_dummies], axis = 1)
df_numeric.head(3)
df_categorical.drop(columns="Species", inplace=True)
```

For the month, we will use cyclic encoding.

```
In [13]: base_angle = 2*np.pi/12
print(base_angle)

cos = np.cos(df_numeric["Month"]*base_angle)
sin = np.sin(df_numeric["Month"]*base_angle)

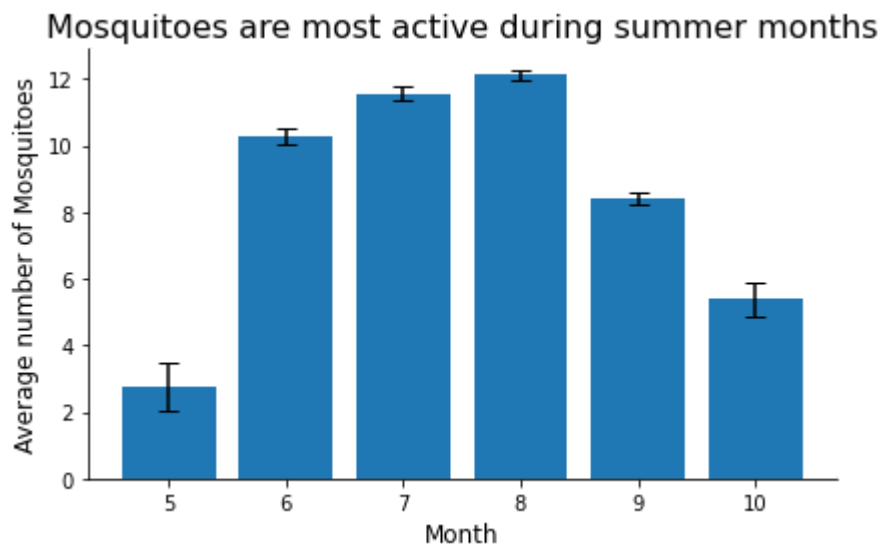
df_numeric["Month_cos"] = cos
df_numeric["Month_sin"] = sin
```

```
0.5235987755982988
```

2. What is the average number of mosquitoes for each month? What trends do you notice?

```
In [14]: monthly_averages = df_numeric.groupby("Month").mean()
monthly_sem = df_numeric.groupby("Month").sem()

plt.figure()
plt.bar(monthly_averages.index.values, monthly_averages["Mosquito number"], yerr=monthly_
plt.xlabel("Month", size=12)
plt.ylabel("Average number of Mosquitoes", size=12)
plt.title("Mosquitoes are most active during summer months", size=16)
sns.despine()
plt.tight_layout()
plt.show()
```




Part 2 - Statistical Analysis

1. Is there a statistically significant difference between the different mosquito species when looking at the occurrence of West Nile Virus?

```
In [15]: species_averages = mosquito_df.groupby("Species")["WNV Present"].aggregate(["mean", "std"])

plt.figure()
plt.bar(species_averages.index.values, species_averages["mean"]*100, yerr=species_averages["std"]*100)
plt.xlabel("Species", size=12)
plt.ylabel("WNV Present (%)", size=12)
plt.title("Mosquito species strongly influences WNV Presence", size=16)
sns.despine()
plt.tight_layout()
plt.show()
```



I will first test to see if there is any significance between the populations using an ANOVA test.

```
In [16]: anova_data = {}
mosquitoes = mosquito_df["Species"].unique()
```

```

for species in mosquitoes:
    anova_data[species] = mosquito_df.loc[mosquito_df["Species"] == species, "WNV Prese

stats.f_oneway(anova_data["CULEX PIPIENS"],
               anova_data["CULEX RESTUANS"],
               anova_data["CULEX SALINARIUS"],
               anova_data["CULEX TERRITANS"])

```

Out[16]: F_onewayResult(statistic=105.45270503888439, pvalue=1.082615440825039e-67)

Results from ANOVA indicate that there is at least one species which is significantly different from the others. I will now run a proportions z test between each species pair individually.

```

In [17]: mosquitoes = mosquito_df["Species"].unique()

sums = []
for key in anova_data:
    sums.append(anova_data[key].sum())

counts = []
for key in anova_data:
    counts.append(anova_data[key].count())

for n in range(4):
    for m in range(n+1, 4):
        compare_sums = []
        compare_sums.append(sums[n])
        compare_sums.append(sums[m])

        compare_counts = []
        compare_counts.append(counts[n])
        compare_counts.append(counts[m])

        print("p value (" + mosquitoes[n] + " vs " + mosquitoes[m] + "):")
        print(proportions_ztest(compare_sums, compare_counts)[1])
        print("\n")

```

p value (CULEX RESTUANS vs CULEX TERRITANS):
4.149911157085309e-48

p value (CULEX RESTUANS vs CULEX SALINARIUS):
1.2461111009600203e-09

p value (CULEX RESTUANS vs CULEX PIPIENS):
1.0019767964463832e-09

p value (CULEX TERRITANS vs CULEX SALINARIUS):
0.0021029008936009263

p value (CULEX TERRITANS vs CULEX PIPIENS):
8.231567115719057e-60

p value (CULEX SALINARIUS vs CULEX PIPIENS):
1.164688112052138e-12

The proportions z tests allow me to conclude that there is a significant difference between all species.

2. Which columns are positively correlated with the number of mosquitoes caught? Which columns are negatively correlated? Are these correlations statistically significant?

I will determine the linear correlation between `Mosquito number` and all other variables using the `corr()` function.

```
In [18]: correlation = df_numeric.corr()
correlation = correlation["Mosquito number"]
correlation.pop("Mosquito number")
correlation = correlation.sort_values(ascending=False)
correlation
```

```
Out[18]: WNV Present      0.408034
Year      0.129326
Trap type_SENTINEL  0.108575
Lat      0.096820
Species_CULEX RESTUANS  0.070999
Species_CULEX PIPIENS  0.014730
Month_sin  0.005443
Trap type_OVI -0.005392
Month     -0.040426
Month_cos -0.064980
Trap type_GRAVID -0.138275
Species_CULEX TERRITANS -0.150962
Lon      -0.151421
Name: Mosquito number, dtype: float64
```

The following columns have a positive correlation with `Mosquito number` :

- WNV Present
- Year
- Trap type_SENTINEL
- Lat
- Species_CULEX RESTUANS
- Species_CULEX PIPIENS
- Month_sin

The following columns have a negative correlation with `Mosquito number` :

- Trap type_OVI
- Month
- Month_cos
- Trap type_GRAVID
- Species_CULEX TERRITANS
- Lon

The strongest correlation is with `WNV Present`, which only has a correlation value of about 0.41, indicating that none of these values has a strong linear correlation with `Mosquito number`. It should be noted that in this experiment, we are only looking at a linear correlation. Trying to correlate these data using other function types, such as logarithms, may result in stronger correlation values.

Part 3 - Advanced Statistical Analysis

1. Run a linear regression to determine how the independent variables affect the number of mosquitoes caught. Explain your model construction process. Analyze the model and the results, and discuss the model's limitations. This may end up being an iterative process.


Note:

- You will likely see a low R^2 value, that is to be expected.
- This dataset does not respond well to performing VIF analysis, so this is not required.
- `WNV Present` **must not** be one of your independent variables.

```
In [20]: x_vals = df_numeric[["Year", "Month", "Lat", "Lon", "Month_cos", "Month_sin"]]
         y_vals = df_numeric["Mosquito number"]
```

```
In [21]: df_corr = x_vals.corr()

plt.figure(figsize=(6,6))
sns.heatmap(df_corr.round(2), vmin=-1, vmax=1, cmap="coolwarm", annot=True, mask=np.tr
plt.show()
```




```
In [22]: x_vals_const = sm.add_constant(x_vals)

mosquito_lr = sm.OLS(y_vals, x_vals_const)
mosquito_lr = mosquito_lr.fit()
display(mosquito_lr.summary())
```

OLS Regression Results						
Dep. Variable:	Mosquito number			R-squared:	0.047	
Model:	OLS			Adj. R-squared:	0.047	
Method:	Least Squares			F-statistic:	152.5	
Date:	Mon, 09 Aug 2021			Prob (F-statistic):	1.10e-189	
Time:	17:11:21			Log-Likelihood:	-73899.	
No. Observations:	18495			AIC:	1.478e+05	
Df Residuals:	18488			BIC:	1.479e+05	
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-2715.6779	111.802	-24.290	0.000	-2934.820	-2496.536
Year	0.4063	0.026	15.473	0.000	0.355	0.458
Month	-0.2387	1.540	-0.155	0.877	-3.258	2.780
Lat	-4.9906	1.267	-3.940	0.000	-7.474	-2.508
Lon	-24.0731	1.533	-15.706	0.000	-27.078	-21.069
Month_cos	-6.4370	2.518	-2.556	0.011	-11.372	-1.502
Month_sin	-6.5608	2.182	-3.007	0.003	-10.837	-2.284
Omnibus:	5024.299	Durbin-Watson:	1.538			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	10667.030			
Skew:	1.622	Prob(JB):	0.00			
Kurtosis:	4.820	Cond. No.	2.33e+06			


Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.33e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [23]: # soft predictions, the probability of getting the deposit
y_pred = mosquito_lr.predict(x_vals_const)

plt.subplots(2, figsize=(8,5))
```

```
plt.subplot(1,2,1)
plt.scatter(y_vals, y_pred)
plt.xlabel("Measured Values")
plt.ylabel("Predicted Values")
plt.title("Measured vs. Predicted")
plt.subplot(1,2,2)
plt.scatter(mosquito_lr.resid, mosquito_lr.fittedvalues)
plt.xlabel("Measured Values")
plt.ylabel("Residuals")
plt.title("Residual Plot")
sns.despine()
plt.show()
```



In [24]:

```
x_vals = df_numeric[["Year", "Lon", "Month_cos"]]
y_vals = df_numeric["Mosquito number"]

x_vals_const = sm.add_constant(x_vals)
mosquito_lr = sm.OLS(y_vals, x_vals_const)
mosquito_lr = mosquito_lr.fit()
display(mosquito_lr.summary())

plt.subplots(2, figsize=(8,5))
plt.subplot(1,2,1)
plt.scatter(y_vals, y_pred)
plt.xlabel("Measured Values")
plt.ylabel("Predicted Values")
plt.title("Measured vs. Predicted")
plt.subplot(1,2,2)
plt.scatter(mosquito_lr.resid, mosquito_lr.fittedvalues)
plt.xlabel("Measured Values")
plt.ylabel("Residuals")
plt.title("Residual Plot")
sns.despine()
plt.show()
```

OLS Regression Results

Dep. Variable: Mosquito number

R-squared:

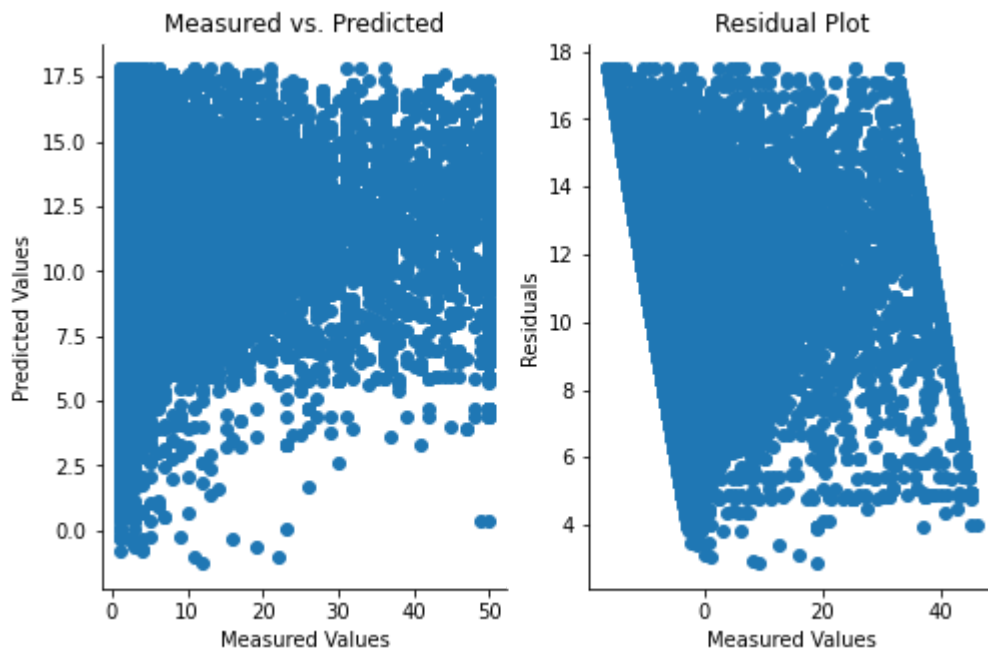
0.039

Model:	OLS	Adj. R-squared:	0.039
Method:	Least Squares	F-statistic:	249.5
Date:	Mon, 09 Aug 2021	Prob (F-statistic):	9.31e-159
Time:	17:11:21	Log-Likelihood:	-73978.
No. Observations:	18495	AIC:	1.480e+05
Df Residuals:	18491	BIC:	1.480e+05
Df Model:	3		
Covariance Type:	nonrobust		
	coef	std err	t P> t [0.025 0.975]
const	-2533.2062	99.826	-25.376 0.000 -2728.873 -2337.539
Year	0.4010	0.026	15.235 0.000 0.349 0.453
Lon	-19.7895	1.023	-19.348 0.000 -21.794 -17.785
Month_cos	-1.9387	0.272	-7.126 0.000 -2.472 -1.405
Omnibus:	5098.704	Durbin-Watson:	1.523
Prob(Omnibus):	0.000	Jarque-Bera (JB):	10941.515
Skew:	1.640	Prob(JB):	0.00
Kurtosis:	4.854	Cond. No.	2.07e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.07e+06. This might indicate that there are strong multicollinearity or other numerical problems.



2. Run a logistic regression to determine how the independent variables affect West Nile Virus presence. Explain your model construction process. Analyze the model and the results, and discuss the model's limitations. This may end up being an iterative process.


Note: `Mosquito number` should be one of your independent variables.

In [25]:

```
x_vals_logis = df_numeric.drop(columns=["WNV Present", "Month"])
x_vals_logis_const = sm.add_constant(x_vals_logis)
y_vals_logis = df_numeric["WNV Present"]

df_corr_logis = x_vals_logis.corr()

plt.figure(figsize=(6,6))
sns.heatmap(df_corr_logis.round(2), vmin=-1, vmax=1, cmap="coolwarm", annot=True, mask
plt.show()
```



```
In [26]: WNV_logistic_r = sm.Logit(y_vals_logis, x_vals_logis_const)
WNV_logistic_r = WNV_logistic_r.fit()
display(WNV_logistic_r.summary())
```

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.370386
Iterations: 35

C:\Users\Daniel\anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to "

Logit Regression Results

Dep. Variable:	WNV Present	No. Observations:	18495
Model:	Logit	Df Residuals:	18482
Method:	MLE	Df Model:	12
Date:	Mon, 09 Aug 2021	Pseudo R-squ.:	0.2901
Time:	17:11:22	Log-Likelihood:	-6850.3
converged:	False	LL-Null:	-9649.5
Covariance Type:	nonrobust	LLR p-value:	0.000
	coef	std err	z P> z [0.025 0.975]
const	-574.0507	26.244	-21.874 0.000 -625.488 -522.614

Year	0.1237	0.007	17.707	0.000	0.110	0.137
Mosquito number	0.0680	0.002	43.020	0.000	0.065	0.071
Lat	-0.7723	0.297	-2.604	0.009	-1.353	-0.191
Lon	-3.9552	0.349	-11.334	0.000	-4.639	-3.271
Trap type_GRAVID	0.2650	0.138	1.926	0.054	-0.005	0.535
Trap type_OVI	-11.1220	6682.078	-0.002	0.999	-1.31e+04	1.31e+04
Trap type_SENTINEL	-0.2750	0.223	-1.231	0.218	-0.713	0.163
Species_CULEX PIPIENS	1.0406	0.327	3.181	0.001	0.400	1.682
Species_CULEX RESTUANS	0.8890	0.327	2.722	0.006	0.249	1.529
Species_CULEX TERRITANS	-1.3407	0.424	-3.163	0.002	-2.171	-0.510
Month_cos	-3.3889	0.149	-22.700	0.000	-3.681	-3.096
Month_sin	-7.0518	0.218	-32.282	0.000	-7.480	-6.624

In [27]:

```
# soft predictions, the probability of getting the deposit
y_proba = WNV_logistic_r.predict(x_vals_logis_const)

y_pred = np.where(y_proba >= 0.5, 1, 0)

num_correct = (y_pred == y_vals_logis).sum()

accuracy = num_correct / y_pred.shape[0] * 100

print(f"The model accuracy is {round(accuracy,2)}%")
```

The model accuracy is 83.04%

Model	Accuracy
Model v1	83.04%

All trap types have p values above 5%. I will try dropping all of them.

In [28]:

```
x_vals_logis = df_numeric.drop(columns=["WNV Present", "Month", "Trap type_GRAVID", "Tr
x_vals_logis_const = sm.add_constant(x_vals_logis)
y_vals_logis = df_numeric["WNV Present"]

WNV_logistic_r = sm.Logit(y_vals_logis, x_vals_logis_const)
WNV_logistic_r = WNV_logistic_r.fit()

y_proba = WNV_logistic_r.predict(x_vals_logis_const)
y_pred = np.where(y_proba >= 0.5, 1, 0)
num_correct = (y_pred == y_vals_logis).sum()
accuracy = num_correct / y_pred.shape[0] * 100
print(f"The model accuracy is {round(accuracy,2)}%")
```

Optimization terminated successfully.

Current function value: 0.370737

Iterations 9

The model accuracy is 82.92%

Model	Accuracy
-------	----------

Model	Accuracy
Model v1	83.04%
Model v2	82.92%

Slight loss in accuracy when trap type is removed. Trap type will stay in the model. I will also try dropping longitude and latitude, seperately.

In [29]:

```
x_vals_logis = df_numeric.drop(columns=["WNV Present", "Month", "Lon"])
x_vals_logis_const = sm.add_constant(x_vals_logis)
y_vals_logis = df_numeric["WNV Present"]

WNV_logistic_r = sm.Logit(y_vals_logis, x_vals_logis_const)
WNV_logistic_r = WNV_logistic_r.fit()

y_proba = WNV_logistic_r.predict(x_vals_logis_const)
y_pred = np.where(y_proba >= 0.5, 1, 0)
num_correct = (y_pred == y_vals_logis).sum()
accuracy = num_correct / y_pred.shape[0] * 100
print(f"The model accuracy w/o Lon is {round(accuracy,2)}%")

x_vals_logis = df_numeric.drop(columns=["WNV Present", "Month", "Lat"])
x_vals_logis_const = sm.add_constant(x_vals_logis)
y_vals_logis = df_numeric["WNV Present"]

WNV_logistic_r = sm.Logit(y_vals_logis, x_vals_logis_const)
WNV_logistic_r = WNV_logistic_r.fit()

y_proba = WNV_logistic_r.predict(x_vals_logis_const)
y_pred = np.where(y_proba >= 0.5, 1, 0)
num_correct = (y_pred == y_vals_logis).sum()
accuracy = num_correct / y_pred.shape[0] * 100
print(f"The model accuracy w/o Lat is {round(accuracy,2)}%")
```

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.373923
Iterations: 35

The model accuracy w/o Lon is 82.68%

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.370570
Iterations: 35

The model accuracy w/o Lat is 82.95%

C:\Users\Daniel\anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
C:\Users\Daniel\anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")

Model	Accuracy
Model v1	83.04%
Model v2	82.92%
Model v3a	82.68%
Model v3b	82.95%

Model v1 is the most accurate.

